

Lecture 1

1/4/2016

Red Star 3.0 (from fedora)

- automatic watermarking (steganography)
- tamper-resistant OS modification

<http://web.cs.ucla.edu/classes/winter16/cs111>

Course organization and grading

17 lectures

1/9 1 midterm (during lecture) 100 minutes

2/9 1 final exam 180 minutes

1/3 4 labs (teams of 2) 1/2 each

shell, kernel hacking, filesystem, networking

miscellaneous

2/15 individual minilabs 1/15 each

scheduling, virtual memory with partner

1/12 1 design problem (lab extension), requires written report

1/15 1 2-3 page paper on an operating system topic

1/20 scribe notes (groups of up to 4), webpage for lecture
Due week after lecture

Open book / notes / assignments on the exams :-)

What's a System?


OED original (1928)

I. An organized or connected group of objects

II. A set of principles, etc., a scheme, method
from Greek $\sigma\tau\epsilon\rho\gamma\mu\alpha$ - organized whole,
government, constitution, a body of people or animals,
musical interval


roots:

- set up with "



Book's definition:

- A set of interconnected components that has a specified behavior observed at the interface of its environment.



Example of subsystems,
not seen from environment

Operating System - American Heritage Dict. (2000)

Software designed to control hardware of a specific data processing system in order to allow users and application programs to make use of it.

Encarta (2007).

master control program in a computer

Wikipedia - v698216816 (2016-01-04)

collection of smaller programs and software that is used to control and operate the computer system

Some major missing issues...

- resource management
- reliability + error handling
- security

circa 2008 hardware

\$ls -l big

$\sim 10^{19}$

-rw-rw-r-- 1 'eggert faculty 922337203685471500

-Oct 6 11:31 big ← ZFS

Zeta File System

\$grep x big

10^{21} bytes/second

\$time grep x big

real 0m0.009s 10^{-2} s

10^{22} bits/second

<http://what-if.xkcd.com/31/> 167

Internet bandwidth ~~167~~ Tb/s $\sim 10^{11}$ bits/second

This means it's possible to physically move data faster than it is to send it through the Internet, given a sufficient size

\$1.2 million/gallon → all freight trucks in the U.S.

~ 0.52 $\frac{\text{bits}}{\text{bytes/second}}$

ZFS has intensional files - the data is not stored on the disk, rather, it is a description or a program that represents the data.

= It also has extensional files, which are interpreted directly.

The file above was created with...

\$truncate 922... 000 f

So how does grep work so quickly?

If it is an international file, then it just reads that.
Hence, in this case grep quickly realizes that there is no x and returns almost immediately.

Aside...

This is useful for virus scanners to exploit. If wants to skip parts that cannot possibly be viruses. That way, it catches them faster.

Problem Areas in OS Design

- Incommensurate scaling

not everything scales at the same rate

cost/unit ↓ economies of scale (Adam Smith, pin-factory)

cost/unit ↑ diseconomies of scale (star network)

Things break as you grow with diseconomies.

Economies can cause waste.

- Emergent properties

larger systems have properties that smaller ones didn't

• (Tacoma Narrows bridge) ← resonance frequency

• UCLA campus network → Napster


- Propagation of effects

pathways for propagation are "more effective"

in a digital system, one component change can be very profound

international characters

⌘ ← 13091257



file ← C:\Windows\foo.txt should become ⌘.txt

Change Δ breaks part X.

Turns out, the coding was shift-JIS

where first byte is all 1's, and the second could be anything.

However: the 2nd byte could be the ASCII '\', and the file system will not work even though it wasn't changed.

SOLUTION: use UTF-8.

- Trade Offs

Waterbed Effect - if we push one part of the waterbed we must send the water somewhere else.

[ex] if we used the SOLUTION, we would need to represent Japanese characters with base 3 characters.

[ex] Too easy to break into UCLA registrar passwords too easy to guess

solution: keycards.

Tradeoff: extra work for maintenance

Benefit: Security.

Lecture 2

1/6/2016

Today's topics:


- a bit more philosophy

- how not to make an OS

Complexity

that are economically viable

Moore's Law - # of transistors doubles approx every year




Log scale! bits/chip

However, in recent years the exponential growth has begun to trail off.

Date of

Introduction

Kryder's Law secondary storage capacity



However, complexity is increasing at a higher rate

Why Exponential Growth?

use to
design the
UNIVAC II

→ UNIVAC I (slow, safe, hand designed)

UNIVAC II

$$\frac{d(\text{technology})}{dt} = K * \text{technology}, \text{ which comes out to be exponential.}$$

How Not to Make an OS

First off, why shouldn't we make an OS?

- simplicity
- performance
 - speed - memory
- Reliability (as a consequence of simplicity)
- security (minimize code I must run.

Consider that I am submitting a proposal for a business on Friday and I want to make my own system (I'm paranoid)

no OS
Application - count of words in an ASCII text file,

Interface

(bytes: '\001' - '\177')
'1' - '\127'

~10 year old desktop

① Power switch

BIO's - where is program, and where is text file?

Core i3-4160 (3 MiB cache, 4 GiB dual channel DDR3

file is here → 1 TB hard drive, SATA, 7200rpm in-adjacent sectors Intel 4400 graphics

39716

Traditional disk drive, 512 bytes each sector


Cycling power clears CPU, cache, and RAM

Bootstrapping Problem - if we know ..., then, but where do we start? We need to have something to work with in order to get started

built into hardware

instruction pointer - $0xffff0$ $2^{20}-16$

access a special region, $0xfff0$



ROM region ← hardwired by manufacturer to do what we want it to do.

-or-
PROM
(programmable read-only memory)

we program, if we got the manufacturer to produce it for us. Most paranoid approach

-or-


called "nonvolatile memory"

EEPROM

(electronically erasable PROM)

Downsides - might render computer unbootable if bad erasure. Erasure takes time and effort

We want the word count program to be on disk, and use the EEPROM to get to it somehow. The problem is, we need the program in RAM so the instruction pointer can use it.



location of wc program, size of wc + loading program

Problem: * still

program dependent.

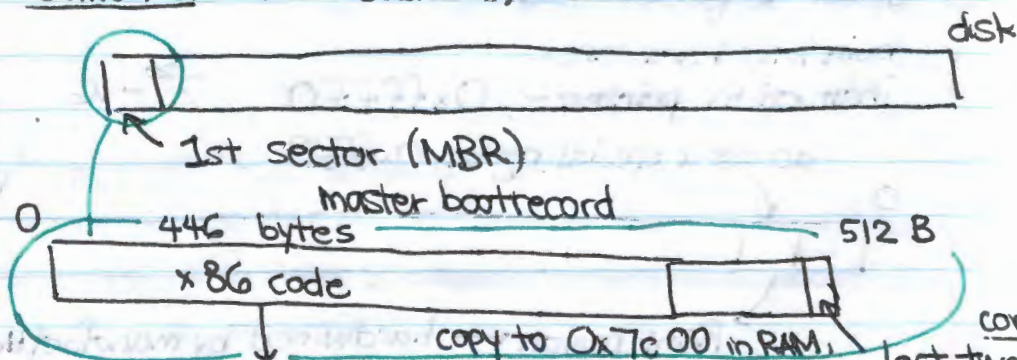
We don't want the

EEPROM to have to

change.

and so the solution is to have the size and location to be the same for every program in EEPROM.

Convention: (It's bad :))




• hardware sanity checks, then jump in.

• some CPU checks - checks RAM


• checks for devices in order

finds 1st device with an MBR

convention:
Ox 55 Ox AA
little endian
Ox AA 55



Firmware (EEPROM) → MBR (first sector on disk) →
do whatever to code so long as it is in first 446 bytes



The controller communicates through the buses via their registers

0x1f7 → status register 1B contains info about state of disk's controller

0/1

controller is available!

~~1f7 status, command~~
~~1f2 sector count~~
 1f0 read data
 1f2 sector count
 1f3 — low order byte
 1f4
 1f5
 1f6 — high order byte
 1f7 status, cmd

Sector #

```

* while ((inb(0x1f7) & 0x052) != 0x40)
    continue;
    insl(0x1f0, a, 128);
  }
  
```

← sizeof(int)

✓ We can have a copy in firmware at some location

-or-

✗ it can be in the master boot record

-or-

✗ it can be in the word count program

Isn't this wasteful? Let's just use one copy

BIOS = basic input/output system

downside: this must be built into the system

V₂

Final O

```

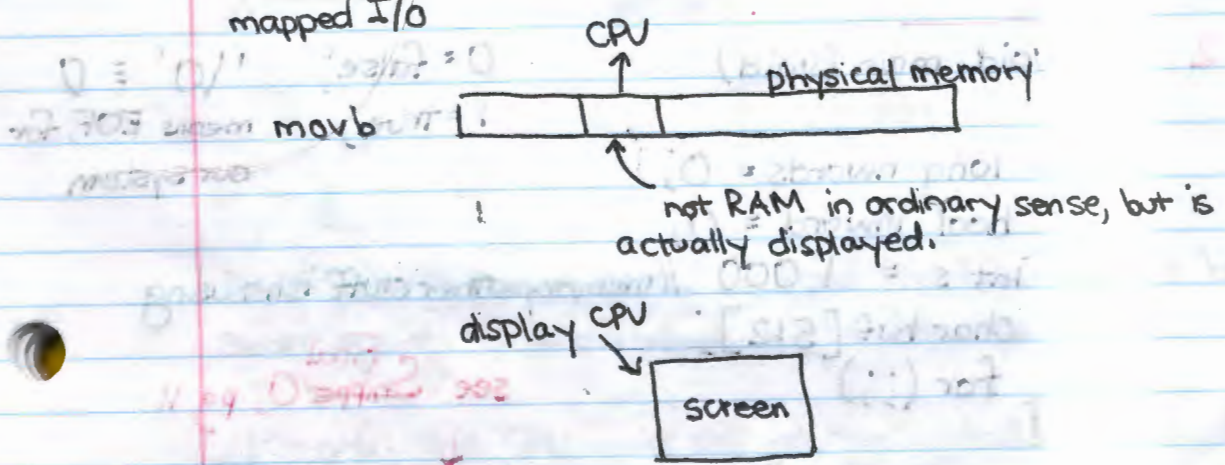
void read_id_sector (int s, char* a)
{
    while ((inb(0x1f7) & 0xc0) != 0x40)
    {
        continue;
        // While not free, continue;
        outb(0x1f2, 1); // Tells that sector count is just 1
        outb(0x1f3, s); // Write low 8 bits to s.
        outb(0x1f4, s >> 8); // Write next 8 bits to proper spot in s.
        outb(0x1f5, s >> 16); // Write next 8 bits to proper spot in s.
        outb(0x1f6, s >> 24); // Write next 8 bits to proper spot in s.
        outb(0x1f7, 0x20); // change status to reading the sector
    }
    while ((inb(0x1f7) & 0xc0) != 0x40)
    {
        continue;
        insl(0x1f0, a, 128); // Read 128 longs worth of memory into RAM
    }
    // starting from the mem address
    512 / sizeof(int)
}
    
```

(status & 11000000) != 01000000
 01
 sector is free

Lecture 3

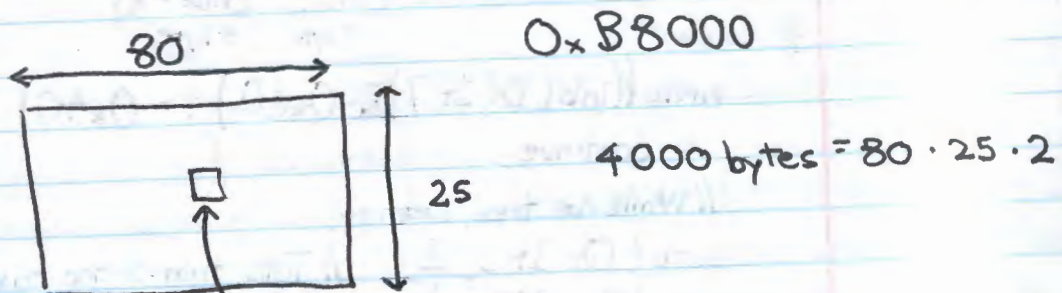
1/11/16

To output to screen, not programmed I/O (PIO) but memory mapped I/O



Standard:

Model Screen as 80x25 grid of characters



16 bits to represent
low order 8 bits is ASCII

high order bits = graphical appearance

V1 cleaner in v2 pg

Snippet 1

```

void output_to_screen (long a)
{
    short long *p = (long*) 0xB8000 / 2 + (80 * 25 - 80) / 2;
    // Won't work! Output will correspond to early bytes on screen.
    do while (n != 0)
    {
        *p++ = (7 << 8) | (n % 10 + '0');
        n /= 10;
    } while (n != 0);
}

```

figure out where on screen to start

reversed order!

otherwise blank screen if 0.

graphic setting

V1

Snippet 2

```

void main(void)
{

```

long nwords = 0;

bool inword = 0;

int s = 1,000 // memory other stuff isn't using

char buf[512];

for (i; i

}

0 = false;

1 = true;

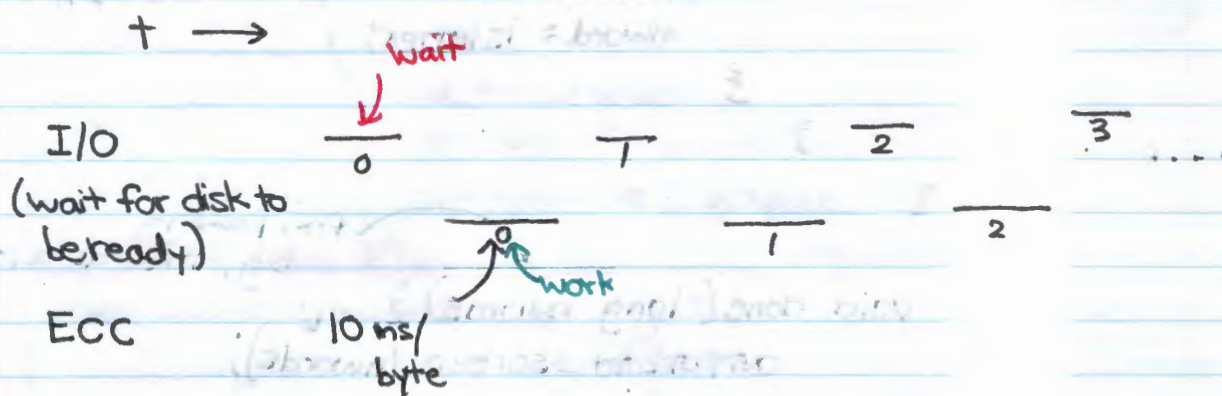
'\0' = 0

means EOF, for our system


Final see Snippet 0, pg 11

This adds complexity! Furthermore, since we're looking at the data in RAM anyway, it's sent to the CPU anyway, so DMA isn't as useful.


However, consider a crypto app; very CPU intensive



How do we improve this? Double buffering, where we load one buffer while working on the other one.



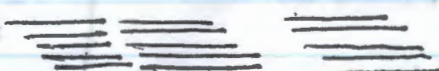
However, if I/O and computation is too high in discrepancy using double buffering is not as worthwhile.



Triple buffering is good when we load; process 1, process 2, etc.!

Multitasking (with multiple processors)

I/O




ECC

Load multiple from I/O
then process one at a time
While loading next set, if
processing is faster by a
significant margin.

How to scale up these programs.


- Fancier performance tricks
without rewriting applications
- Multitasking
without rewriting applications

Modularity



cost of maintaining a module
of N lines of code is $O(N^2)$

Abstraction



Finding natural divisions in the
programs that makes the program
is easier to manage.

How do you measure the quality of modularity + abstraction?

Simplicity (ease of use)
(ease of learning)

Robustness (tolerance of errors, large inputs) harsh conds
Performance modularity costs this, minor costs unavoidable,
avoid major costs as often as possible.

Flexibility/Lack of Assumptions/Neutrality


Consider....

char *readline (FILE *f);

BAD DESIGN for O.S.

Why?

- ① Performance × (unbounded work)
(unbatched work) ← overhead for small lines
- ② Robustness × (apps crash for big lines)
- ③ Neutrality × (forces particular line ending convention on the app)
- ④ Simplicity ✓ 😊



Let's now examine read-ide-sector...

improve
robust-
ness by
reporting
potential
error
msgs.

void read-ide-sector (int s, char *buf)
→ int read-ide-sector (int s, char *buf)
int read (int byte-offset, char *buf, int bufsize)

↑
more general, no need to worry about exact sector size

int read (int byte-offset, void *buf, int bufsize)

↑
read other things besides chars

off_t
int read (int byte-offset, void *buf, size_t bufsize)

// long long on x86-64

↑
increase max bufsize, make
this portable

// size_t is unsigned long on x86-64

ssize_t read (off_t byte-offset, void *buf, size_t bufsize)

↑
bytes read
or -1 if error

Other improvements....

have it return error code? Not the
convention, sadly.

- which file? (opaque file handle) where did the byte-offset go?
- ↓
- `ssize_t read (int fd, void *buf, size_t bufsize);`

BIG IDEA OF UNIX

everything outside the program is a file
(disk, flash drive, network, mouse, keyboard, display)

random access
device

stream device



doesn't need/use the offset, hence its removal

but we lost random access! How do we get it back?

- `size_t lseek (int fd, off_t where value, int flag);`

Corollary, the OS records current file positions starts at 0, after reading n we add n .

- `pread ...`
add complexity to get performance


Mechanisms for Modularity

- Function calls
 - + simple and well understood
 - + reasonably fast
 - things can go wrong
- Client-server (works, but....slow)
- Virtualization

Lecture 4: OS Organization

1/13/16

Bootstrapping UEFI



GUID: globally unique ID for disk partitions (128-bit)

Standardized - GUID partition table (GPT)

Standard layout in each partition:

FAT 12

kernel with known name

FAT 16

FAT 32

Firmware is in charge until booting. This is enforced organization.

How to enforce modularity after booting
function calls (terrible way?)

```

char buf[2000];
read(3, buf, 1000);

int fact(int n) {
    if (!n) return 1;
    return n * fact(n-1);
}

// unoptimized:
pushq %rbp
movq %rsp, %rbp
subq $16, %rsp
movl $0, -4(%rbp)
cmpl $0, -4(%rbp)
jne .L2
movl $1, %eax
jmp .L3
  
```

Recall, %edi is parameter
%rax is return value

.L2

```

movl    -4(%rbp), %eax
subl    $1, %eax
movl    %eax, %edi
call    fact
imull    -4(%rbp), %eax

```

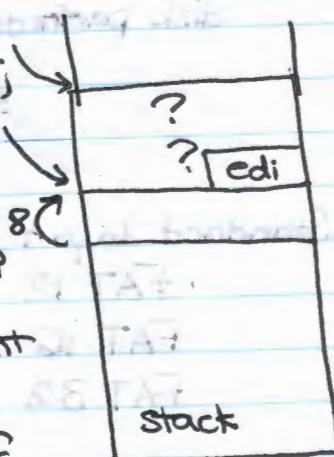
.L3

POP

```

leave   rsp=rbp; rbp=
ret     rip=*rsp++;

```



From start to finish:

- Decrement by 8
- Set `%rbp` to `%rsp`'s current location
- Move up 16 bytes, 2 sets of
- Junk inside.
- low order 4 bytes of `%edi`

is set to where `%rbp`'s

- Perform comparison, if `%edi` is 0, go to L3, store 1 in `%eax` (low 4 bytes of `%rax`), then return

- Otherwise we jump to L2, first two commands get `n-1` into `%eax`.

```

movl    %eax, %edi moves %edi

```

- Then we call factorial

- We then multiply into our `%eax` register which will contain the answer.

Things that can go wrong:


fact (INT_MIN) ← overflow, attempt fact (INT_MAX)

fact (50) ← will only give us the bottom 32 bits (4 bytes) of the correct answer

However, each time we call it we will put up 12 bytes, eventually obliterating the rest of the data structures (or dump core, if we're lucky).

We can catch this but they have some downsides...

- Adding checks will greatly slow it down
- Guard page



← if it reaches here we will trap

Suppose in the earlier machine code we call fact2 instead of fact. What can fact2 do to mess up fact?

- Modify stored values in fact (activation record)
 - make it return to the wrong spot
 - modify the return value
 - Cause an infinite loop
 - Jump elsewhere in the program

Can fact mess up fact2?

Put a random value into %edi

- movl \$0, %esp ← will have fact2 enter forbidden zone as soon as it tries to push to the stack
- jump fact2

We have soft modularity.

- it doesn't scale to large applications

We want hard modularity, where a failure in a module won't tank the rest of them.

- virtualization Unidirectional hard modularity
- client-server Hard modularity in both directions

Simplest way to get virtualization

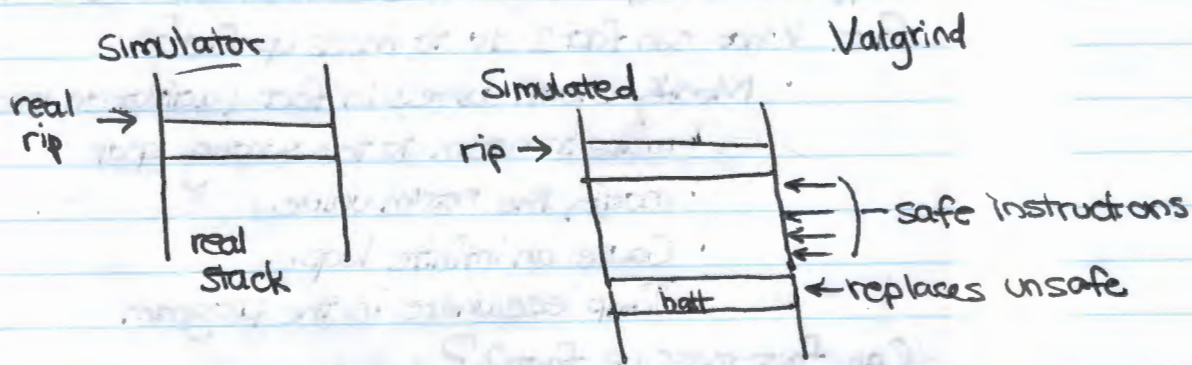
write a simulator for the machine, the app runs on
carefully checks instructions

- halt \rightarrow return
- autoincrement \rightarrow return
- count instructions
- when it reaches limit \rightarrow return

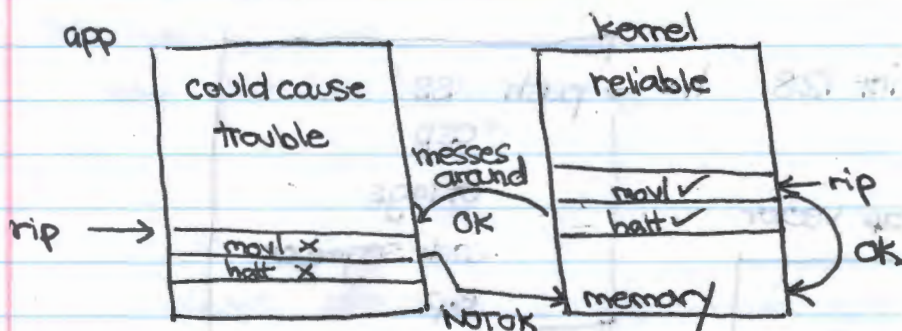
Unfortunately, this method is slow. :-)

In fact, it's usually too slow for production

We need a virtualizable machine, a virtualizable processor



There's one big problem left... it's still too slow!!!



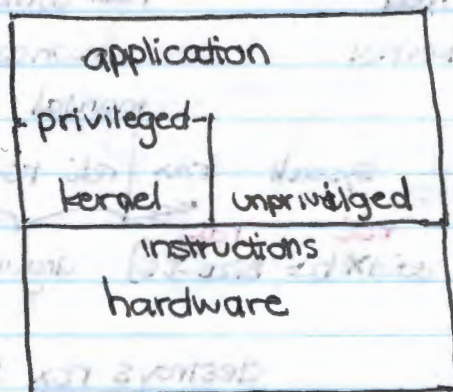
Protected Transfer of Control

should make these rare for this model to work
~~dangerous~~ privileged instructions in the application do not execute.

rip becomes in kernel

We have a privileged bit, which is set to 0 in the app. When it encounters a privileged instruction, it passes to the kernel, sets the privileged bit to 1, and does what it needs to do.

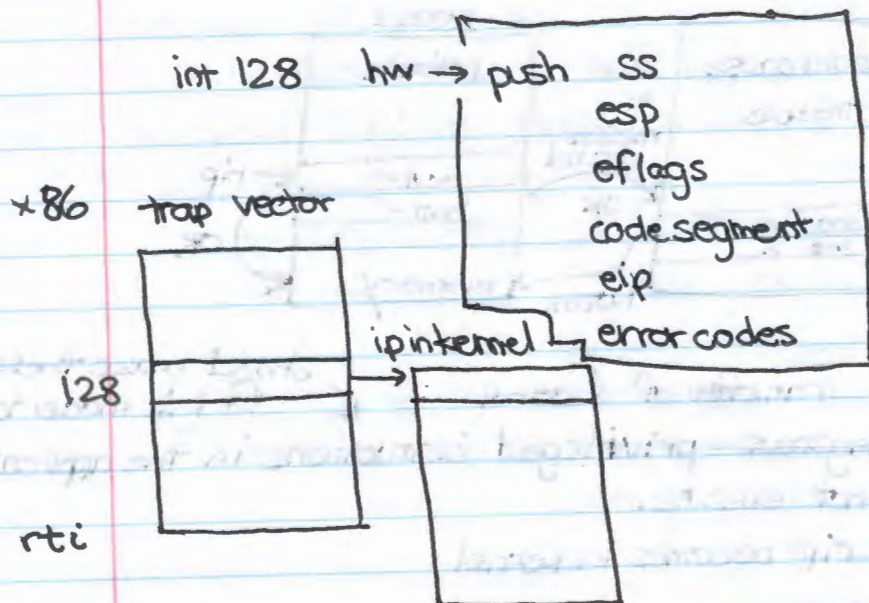
Layered architecture



Classic way to enter the kernel:

execute a privileged instruction with a standard convention

this is \rightarrow int 128



rti: int as ret: call
expensive cheap

Nowadays, for the x86, x86-64, we have a special sys call machine instruction:

- When executed
 - enters kernel

rax contains syscall #'s, which can be found in a manual

syscall rax | rdi rsi rdx r10 r8, r9
rdi rsi rdx arguments

{

asm;
deal with errno;

}

destroys rcx, r11

result into rax

-4095 -1 means failure -errno

So, if rax is wrong then we send result to errno and return -1.

Mechanism works

Now, what does the user see?

↑
app
dev

Model: processes: programs running in a virtual machine atop an operating system

Creation `pid_t fork(void);` Clones current process
returns 0 in child
returns child's pid in parent
returns -1 on failure

Destruction `void _Noreturn _exit(int);`
↑ ↑
never bypass
returns normal
cleanup

Calling process immediately stops running, but not gone!

Process object

exit-status	pid
fd-table	

So this isn't the true way to destroy a process

`pid_t waitpid(pid_t, int*, int)`
status ← flags - \emptyset - NOHANG