

Discussion 2: Signal Handling

Lab 1B

signal handling

✓ -- abort

Get it to segmentation fault!

Do it by referencing a pointer set to NULL;

— catch N.

Catch any processes that send signal N.

```
#include <signal.h>
```

```
#include <stdlib.h>
```

```
void segfault_sigaction (int signal, siginfo_t *s,  
void *arg) { exit(signal); }
```

```
void segfault_sighandler (int signal)  
{ exit(signal); }
```

```
int main
```

```
struct sigaction sa;
```

```
sa.sa_sigaction = segfault_sigaction)
```

```
// sa.sa_handler = segfault_sighandler
```

```
sigaction (SIGSEGV, &sa, NULL)
```

```
int *a = NULL
```

```
int b = *a;
```

exits needed or else an attempt to

be made to reexecute a violating

command will be made forever

-- pause

We need an external segment fault...

How do we achieve this?

```
#include <unistd.h> // for pause
```

20
(probably is + void pauseHandler (int signal, siginfo_t * si, void * arg)

}

// do something

}

int main ()

{

struct sigaction sa;

sa.sa_sigaction = pauseHandler;

sigaction (SIGSEGV, &sa, NULL);

sigaction (SIGINFO, &sa, NULL);

pause ();

printf ("continue");

--ignore(:)

Two ways to try to do it:

- Choose to change the instruction pointer (rip)

rax: 64

eax: 32

ax: 16

ah: 8 → high

al: 8 → low

define _GNU_SOURCE

include <signal.h>

include <ucontext.h>

include <sys/main.h>


```

void sigfault_sigaction(int signal, siginfo_t *si, void *arg)
{
    // when ignore is enabled
    ucontext_t *context = (ucontext_t *) arg;
    context->uc_mcontext.gregs[REG_RIP]++;
}

```

Lecture 5:

1/20/16

Big OS Goals: A Recap

Protection Robustness Flexibility Simplicity

Utilization Performance Abstraction Modularity


x86 supports 4 levels of abstraction



Linux does only 2: apps and kernel


- faster!
- simpler

Here's a different approach:



fork() clone a process, except for ~~process ID~~

- ① process ID \rightarrow pid_t getpid(void);
- ② parent ID \rightarrow pid_t getppid(void); \leftarrow return parent's pid
- ③ file descriptor tables



/dev/urandom is not really random. It uses entropy to seed a RNG.

The child has a file description.

- ④ CPU time info
- ⑤ Pending signals
- ⑥ File locks

↑
execvp(file, argv)
(char const *, char * const*)

It destroys a process, and throws everything away, except:

You destroy:

- all variables
- instruction pointer %rip
- registers
- program

.. We must also reset the signal handlers


```

    bool void printdate(void)
    {
        pid_t p = fork(); // ensure that we don't terminate as soon as
        switch(p)         // the call is over
        {
            case -1: error();
            case 0: { execvp("/bin/date");
                → static char const date[] = "/bin/date";
                execvp (date, (char const *) { date, NULL });
            }
            default: error();
        }

        int status;
        if (waitpid(p, &status, 0) != p) true if process
            error();                      exited
        return (WEXITSTATUS(status) == 0);
    }
}

```

If we were to pass in a parameter `char const *outfile` for this function, we would open it after the child is successfully created (→)

```

    int fd = open(outfile, O_WRONLY);
    if (fd < 0) error();
    if (dup2(fd, 1) < 0) error();
    if (fd != 1) close(0);

    if (fork() == 0)
    {
        ← done before execvp(...)
        // housekeeping
    }
}

```


Another totally different school of thought handles all the housekeeping in one system call.

```
int posix_spawnvp (pid_t * restrict pid,;
```

What does restrict do? It prevents aliasing, allowing
It's a keyword. for additional optimizations

ex: `char * strcpy (char * restrict dest, char * const restrict src)`

Posix_spawnvp is actually more complicated in terms of parameters.


```
int posix_spawnvp ( pid_t * restrict pid,
                    char const * restrict file,
                    posix_spawn_file_actions_t const *
                    restrict file_acts,
                    posix_spawn_attr_t const *
                    restrict attr,
                    char * const * restrict argv,
                    char * const * restrict envp );
```

Look at how annoying this is to use!

Upside: it's fast.

Orthogonality

In math:



I can pick any x_0 , then freely pick any y_0 and z_0 and form the coordinates (x_0, y_0, z_0) .

We want our system calls to be orthogonal to each other.

Orthogonality works with processes

Processes interacting with files, it is slow and unreliable; problems involving robustness and performance.

Network

mouse

keyboard

VS

flash

disk

spontaneous data generation

Infinite*

Stream

random access

request/response

finite

UNIX BIG IDEA

everything is a file.

open lseek

close

read

write

dup2

The principle is to make a set of operations that works with a lot of different things.


One downside is that we may have ops and files that are incompatible with each other.

There's another situation we must cover: when processes communicate with each other.

Pipes are banded buffers

process (A)
write ('xyz')

kernel memory buffer



pointers r and w

first in first out buffers

If the buffer is written, it is added to the buffer


If the buffer is read, it is removed after reading

If the buffer is full and a write is attempted, the buffer will force the writing to wait.

If the buffer is empty and a read is attempted, the buffer will force the reading to wait.

In this way, the pipe can control the flow of data.

However, in this form any reader will get the messages from any writer.



Race conditions - when behavior of a program changes depending on when things get scheduled.

Say we want to implement a sort.

We might need to create a temporary file.

```
int create-temp-file(void) {
    return open("tmp/sorttemp", O_RDWR|O_TRUNC|O_CREAT,
               0600);
}
```

This breaks because two concurrent sorts will result in an overwritten temporary file!

How do we fix this?

`O_EXCL`: will cause it to not work if the file exists already!

Then, we use a RNG to generate a unique filename for us that will hopefully not cause any overlap.

Now we have primitives that work, but cause race conditions depending on the distance between system calls. This is seemingly contradictory to the principle of orthogonality.